

ppmpfg3mw

March 12, 2024

```
[10]: import torch
import torchvision
import torchvision.transforms as transforms
import torch.nn as nn
import torch.nn.functional as F
import torch.optim as optim
import torch.nn.init as init

import matplotlib.pyplot as plt
import numpy as np

device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')
print(device)
```

cuda

## Model

```
[11]: class MLP_CIFR10(nn.Module):
    # Batch norm = True enables the Batch-Normalization
    def __init__(self, input_size=3*32*32, h1_size = 500, h2_size = 250, h3_size=
    ↪ 100, num_classes = 10, batch_norm = False):
        super(MLP_CIFR10, self).__init__();
        self.batch_norm = batch_norm
        # Hidden Layer 1 with ReLU activation
        self.fc1 = nn.Linear(in_features = input_size, out_features = h1_size);
        if batch_norm == True:
            # print("Batch norm1 enabled")
            self.bn1 = nn.BatchNorm1d(h1_size)
        self.relu1 = nn.ReLU()
        # Hidden Layer 2 with ReLU activation
        self.fc2 = nn.Linear(in_features = h1_size, out_features = h2_size);
        if batch_norm == True:
            # print("Batch norm2 enabled")
            self.bn2 = nn.BatchNorm1d(h2_size)
        self.relu2 = nn.ReLU()
        # Hidden Layer 3 with ReLU activation
        self.fc3 = nn.Linear(in_features = h2_size, out_features = h3_size);
```

```

if batch_norm == True:
    # print("Batch norm3 enabled")
    self.bn3 = nn.BatchNorm1d(h3_size)
self.relu3 = nn.ReLU()
# output layer
self.fc_out = nn.Linear(in_features = h3_size, out_features = num_classes);
self.softmax = nn.Softmax(dim = 1); # we need the output to be y0 to y9
→depicting the classes

def forward(self, x):
    x = torch.flatten(x, 1); # flatten all dimensions except 1st dimension
    →which represents batch size

    # Generate the output of Fully connected layer 1
    out = self.fc1(x);
    if self.batch_norm == True:
        out = self.bn1(out)
    out = self.relu1(out);

    # Generate the output of Fully connected layer 2
    out = self.fc2(out);
    if self.batch_norm == True:
        out = self.bn2(out)
    out = self.relu2(out);

    # Generate the output of Fully connected layer 3
    out = self.fc3(out);
    if self.batch_norm == True:
        out = self.bn3(out)
    out = self.relu3(out);

    # Final output layer
    out = self.fc_out(out);
    # softmax has not been applied here since the nn.CrossEntropyLoss() applies
    →LogSoftmax on an input, followed by NLLoss
    # Refernce for the above statemet : https://pytorch.org/docs/stable/
    →generated/torch.nn.CrossEntropyLoss.html#torch.nn.CrossEntropyLoss
    return out

```

## Initialize the model

```
[12]: model = MLP_CIFR10(batch_norm=False).to(device) # Batch Normalization Turned off
```

## Creating Data Loader

```
[13]: # Normalize input data to 0.5 mean and 0.5 SD
```

```

transform = transforms.Compose(
    [transforms.ToTensor(),
     transforms.Normalize((0.5, 0.5, 0.5), (0.5, 0.5, 0.5))])

batch_size = 4

# load Training Dataset
trainset = torchvision.datasets.CIFAR10(root='./data', train=True,
                                         download=True, transform=transform)
trainloader = torch.utils.data.DataLoader(trainset, batch_size=batch_size,
                                           shuffle=True, num_workers=2)

#Load Testing Dataset
testset = torchvision.datasets.CIFAR10(root='./data', train=False,
                                         download=True, transform=transform)
testloader = torch.utils.data.DataLoader(testset, batch_size=batch_size,
                                          shuffle=False, num_workers=2)

# Different classes present in the data set
classes = ('plane', 'car', 'bird', 'cat',
           'deer', 'dog', 'frog', 'horse', 'ship', 'truck')

```

Files already downloaded and verified

Files already downloaded and verified

### Define Loss Function and Optimizer

```

[14]: # Cross entropy loss
criterion = nn.CrossEntropyLoss()
# Stochastic Gradient Descent Optimiser, with learning rate = 0.001 and
    ↪ Momentum = 0.9
optimizer = optim.SGD(model.parameters(), lr=0.001, momentum=0.9)

```

### Cross verify the Images with Labels

```

[15]: dataiter = iter(trainloader)


image, label = next(dataiter); #image size : torch.Size([4, 3, 32, 32]), Batch
    ↪ size = 4

for i in range(batch_size):
    image_to_show = image.T[:, :, :, i] # number of channels, Height of image, width
    ↪ of image, batch size
    plt.subplot(1, batch_size, i+1)
    plt.imshow(image_to_show)
    plt.title(classes[label[i]])

```

WARNING:matplotlib.image:Clipping input data to the valid range for imshow with

RGB data ([0..1] for floats or [0..255] for integers).  
 WARNING:matplotlib.image:Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers).  
 WARNING:matplotlib.image:Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers).  
 WARNING:matplotlib.image:Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers).



## Train the Model

```
[16]: # start training the model
training_error = [] # TO capture the training error
interm_training_acc = [] # to capture intermediate training accuracy
interm_total = 0
interm_correct = 0
total = 0 # Track the total images
correct = 0 # Track the number of matched labels for images

for epoch in range(10):

    running_loss = 0.0;

    for i, data in enumerate(trainloader):
        inputs, labels = data

        inputs = inputs.to(device)
        labels = labels.to(device)

        # reset the gradients to zero
        optimizer.zero_grad()

        # forward pass
        outputs = model(inputs)

        # Loss calculation with respect to Ground Truth labels
        loss = criterion(outputs, labels)
```

```

# Back prop
loss.backward()

# Gradient update
optimizer.step()

# cumulative loss
running_loss += loss.item()

predicted_probability, predicted = torch.max(outputs.data, 1) # For each
↪input image 10 probabilities for classification will be predicted, # from that
↪we need to pick the maximum probability class as our predicted class
total += labels.size(0)
correct += (predicted == labels).sum().item()
interm_total, interm_correct = total, correct
if ((i) % 2000) == 1999: # print every 2000 mini-batches, each mini
↪batch has 4 images
    avg_training_loss = running_loss / 2000;
    print(f'[Epoch : {epoch + 1}, Step : {i + 1:5d}] , Average Training loss:
↪{avg_training_loss}')
    training_error.append(avg_training_loss)
    interm_training_acc.append((100*correct)/total)
    running_loss = 0.0

# Plot the metrics
fig = plt.figure(figsize=(12, 5))
fig.add_subplot(1,2,1)
plt.plot(training_error,color = 'r', label='Training Loss')
plt.xlabel('Mini batch step')
plt.ylabel('loss')

fig.add_subplot(1,2,2)
plt.plot(interm_training_acc,color = 'g', label='Training Accuracy')
plt.xlabel('Mini batch step')
plt.ylabel('Accuracy')
plt.legend()

plt.tight_layout()
plt.show()

print("\nTraining Finished with accuracy : ", ((100 * correct)//total) )

```

```

[Epoch : 1, Step : 2000] , Average Training loss: 2.050745455235243
[Epoch : 1, Step : 4000] , Average Training loss: 1.757174305766821
[Epoch : 1, Step : 6000] , Average Training loss: 1.6824005272388458

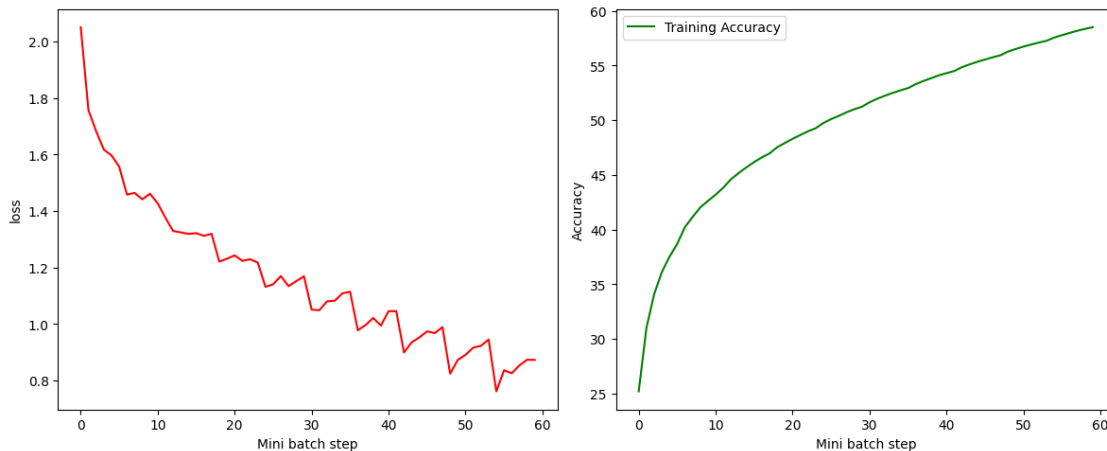
```

[Epoch : 1, Step : 8000] , Average Training loss: 1.6173994615226983  
 [Epoch : 1, Step : 10000] , Average Training loss: 1.5964819515943527  
 [Epoch : 1, Step : 12000] , Average Training loss: 1.5565218421518803  
 [Epoch : 2, Step : 2000] , Average Training loss: 1.4580196150690317  
 [Epoch : 2, Step : 4000] , Average Training loss: 1.4640771860927344  
 [Epoch : 2, Step : 6000] , Average Training loss: 1.4413822045251727  
 [Epoch : 2, Step : 8000] , Average Training loss: 1.4612751085460185  
 [Epoch : 2, Step : 10000] , Average Training loss: 1.427089064463973  
 [Epoch : 2, Step : 12000] , Average Training loss: 1.3762313548624516  
 [Epoch : 3, Step : 2000] , Average Training loss: 1.3292630887776613  
 [Epoch : 3, Step : 4000] , Average Training loss: 1.3241856384351849  
 [Epoch : 3, Step : 6000] , Average Training loss: 1.318869185961783  
 [Epoch : 3, Step : 8000] , Average Training loss: 1.321404395978898  
 [Epoch : 3, Step : 10000] , Average Training loss: 1.3120692287497222  
 [Epoch : 3, Step : 12000] , Average Training loss: 1.319136937595904  
 [Epoch : 4, Step : 2000] , Average Training loss: 1.2207277395948768  
 [Epoch : 4, Step : 4000] , Average Training loss: 1.2309505095407367  
 [Epoch : 4, Step : 6000] , Average Training loss: 1.242813280209899  
 [Epoch : 4, Step : 8000] , Average Training loss: 1.2233950720913709  
 [Epoch : 4, Step : 10000] , Average Training loss: 1.2291937605254353  
 [Epoch : 4, Step : 12000] , Average Training loss: 1.2174601745028049  
 [Epoch : 5, Step : 2000] , Average Training loss: 1.1314555116780103  
 [Epoch : 5, Step : 4000] , Average Training loss: 1.1406325489468871  
 [Epoch : 5, Step : 6000] , Average Training loss: 1.1695521994419396  
 [Epoch : 5, Step : 8000] , Average Training loss: 1.1339136381633579  
 [Epoch : 5, Step : 10000] , Average Training loss: 1.1515361717455088  
 [Epoch : 5, Step : 12000] , Average Training loss: 1.168821255363524  
 [Epoch : 6, Step : 2000] , Average Training loss: 1.0509158919006587  
 [Epoch : 6, Step : 4000] , Average Training loss: 1.0485002759303899  
 [Epoch : 6, Step : 6000] , Average Training loss: 1.0800385706722737  
 [Epoch : 6, Step : 8000] , Average Training loss: 1.0821661712210626  
 [Epoch : 6, Step : 10000] , Average Training loss: 1.1084798576161266  
 [Epoch : 6, Step : 12000] , Average Training loss: 1.1139543134383858  
 [Epoch : 7, Step : 2000] , Average Training loss: 0.977692641383037  
 [Epoch : 7, Step : 4000] , Average Training loss: 0.9960755651630461  
 [Epoch : 7, Step : 6000] , Average Training loss: 1.0210950183402747  
 [Epoch : 7, Step : 8000] , Average Training loss: 0.9944469231776893  
 [Epoch : 7, Step : 10000] , Average Training loss: 1.0452535498891957  
 [Epoch : 7, Step : 12000] , Average Training loss: 1.0454592692144216  
 [Epoch : 8, Step : 2000] , Average Training loss: 0.8990772376107052  
 [Epoch : 8, Step : 4000] , Average Training loss: 0.9350231147864834  
 [Epoch : 8, Step : 6000] , Average Training loss: 0.9523933949491474  
 [Epoch : 8, Step : 8000] , Average Training loss: 0.9738291351981461  
 [Epoch : 8, Step : 10000] , Average Training loss: 0.9676795521322638  
 [Epoch : 8, Step : 12000] , Average Training loss: 0.9888891695505008  
 [Epoch : 9, Step : 2000] , Average Training loss: 0.8239253270998597  
 [Epoch : 9, Step : 4000] , Average Training loss: 0.8730155742289498  
 [Epoch : 9, Step : 6000] , Average Training loss: 0.8908144894224824

```

[Epoch : 9, Step : 8000] , Average Training loss: 0.9163070750022307
[Epoch : 9, Step : 10000] , Average Training loss: 0.9224836231064982
[Epoch : 9, Step : 12000] , Average Training loss: 0.94497442493774
[Epoch : 10, Step : 2000] , Average Training loss: 0.7614482136009029
[Epoch : 10, Step : 4000] , Average Training loss: 0.8360835166203324
[Epoch : 10, Step : 6000] , Average Training loss: 0.8254879862973467
[Epoch : 10, Step : 8000] , Average Training loss: 0.8532491387241753
[Epoch : 10, Step : 10000] , Average Training loss: 0.8734164414613041
[Epoch : 10, Step : 12000] , Average Training loss: 0.8727805498801172

```



Training Finished with accuracy : 58

Accuracy for 10,000 images

```

[17]: correct = 0
total = 0
count = 0
validation_error = []
prediction_accuracy = []
validation_running_loss = 0.0
model.eval()
# since we're not training, we don't need to calculate the gradients for our
→ outputs
with torch.no_grad():
    for i, data in enumerate(testloader):
        count+=1
        images, labels = data # Batch sized images and labels will be loaded
        → here i.e, 4 images and 4 labels are loaded
        images = images.to(device)
        labels = labels.to(device)

```

```

# forward pass
outputs = model(images)

# loss computation
loss = criterion(outputs, labels)

# Accumulate the loss value
validation_running_loss += loss.item()

predicted_probability, predicted = torch.max(outputs.data, 1) # each
↪ input image will have 10 probabilities for classification,
                                                                    # from
↪ that we need to pick the maximum probability class as our predicted class

total += labels.size(0)
correct += (predicted == labels).sum().item()

if ((i+1) % 50) == 0: # print every 50 mini-batches
    avg_validating_loss = validation_running_loss / 500;
    print(f'[Step : {i + 1:5d}] , Average Validation loss:
↪ {avg_validating_loss}')
    validation_error.append(avg_validating_loss)
    validation_running_loss = 0.0

# Plot
plt.figure(figsize=(12, 5))
plt.plot(validation_error, label='Validation Loss')
plt.xlabel('Mini batch step')
plt.ylabel('Validation loss')
plt.legend()

plt.tight_layout()
plt.show()

print(f'\n Accuracy of the network on the {count*batch_size} test images: {100
↪ * correct // total} %')

```


```

[Step :    50] , Average Validation loss: 0.15277633357048034
[Step :   100] , Average Validation loss: 0.1719592792391777
[Step :   150] , Average Validation loss: 0.13672817236185072
[Step :   200] , Average Validation loss: 0.15954202455282213
[Step :   250] , Average Validation loss: 0.12652308634668588
[Step :   300] , Average Validation loss: 0.13593868684768676
[Step :   350] , Average Validation loss: 0.14842946100234986
[Step :   400] , Average Validation loss: 0.14848214885592462
[Step :   450] , Average Validation loss: 0.1456097714006901
[Step :   500] , Average Validation loss: 0.14607786831259728

```



[Step : 550] , Average Validation loss: 0.1589417405128479  
[Step : 600] , Average Validation loss: 0.15315185540914536  
[Step : 650] , Average Validation loss: 0.16348702858388425  
[Step : 700] , Average Validation loss: 0.1494603400528431  
[Step : 750] , Average Validation loss: 0.14306334841251372  
[Step : 800] , Average Validation loss: 0.16430219042301178  
[Step : 850] , Average Validation loss: 0.16250766471028327  
[Step : 900] , Average Validation loss: 0.14951788021251558  
[Step : 950] , Average Validation loss: 0.14629126362502576  
[Step : 1000] , Average Validation loss: 0.14025746634602546  
[Step : 1050] , Average Validation loss: 0.139748472481966  
[Step : 1100] , Average Validation loss: 0.13152092568576335  
[Step : 1150] , Average Validation loss: 0.1288002579510212  
[Step : 1200] , Average Validation loss: 0.13843094052374363  
[Step : 1250] , Average Validation loss: 0.15066212409734725  
[Step : 1300] , Average Validation loss: 0.14961305627226829  
[Step : 1350] , Average Validation loss: 0.15293689513206482  
[Step : 1400] , Average Validation loss: 0.14541964304447175  
[Step : 1450] , Average Validation loss: 0.1410607739686966  
[Step : 1500] , Average Validation loss: 0.15439578494429587  
[Step : 1550] , Average Validation loss: 0.1428820002414286  
[Step : 1600] , Average Validation loss: 0.15549879717826842  
[Step : 1650] , Average Validation loss: 0.15523128631711006  
[Step : 1700] , Average Validation loss: 0.14228341197967528  
[Step : 1750] , Average Validation loss: 0.1611905579417944  
[Step : 1800] , Average Validation loss: 0.15978457593917847  
[Step : 1850] , Average Validation loss: 0.13088434752821923  
[Step : 1900] , Average Validation loss: 0.1490074161142111  
[Step : 1950] , Average Validation loss: 0.15207097494602204  
[Step : 2000] , Average Validation loss: 0.15134687405824662  
[Step : 2050] , Average Validation loss: 0.1695377493798733  
[Step : 2100] , Average Validation loss: 0.15998603767156602  
[Step : 2150] , Average Validation loss: 0.15601812042295932  
[Step : 2200] , Average Validation loss: 0.15380683758854866  
[Step : 2250] , Average Validation loss: 0.14873118564486504  
[Step : 2300] , Average Validation loss: 0.14649967443943024  
[Step : 2350] , Average Validation loss: 0.14631072840094567  
[Step : 2400] , Average Validation loss: 0.13229155695438385  
[Step : 2450] , Average Validation loss: 0.15352160286903382  
[Step : 2500] , Average Validation loss: 0.16686823827028274



Accuracy of the network on the 10000 test images: 53 %

### Saving the Model

```
[18]: PATH = './cifar_net.pth'
      torch.save(model.state_dict(), PATH)
```

### 3. Checking the Ground Truth Values for the Test data


```
[19]: test_data_itr = iter(testloader)
      test_images, labels = next(test_data_itr)
      test_images, labels = next(test_data_itr)
      # print the test images wiht it's ground truth value
      for i in range(test_images.size(0)):
          image = test_images.T[:, :, :, i]
          plt.subplot(1, batch_size, i+1)
          plt.imshow(image)
          plt.title(classes[labels[i]])
```

WARNING:matplotlib.image:Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers).

WARNING:matplotlib.image:Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers).

WARNING:matplotlib.image:Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers).

WARNING:matplotlib.image:Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers).



### Load the Trained Model

```
[20]: saved_model = MLP_CIFR10(batch_norm=False).to(device)
      saved_model.load_state_dict(torch.load(PATH))
```

```
[20]: <All keys matched successfully>
```

```
[21]: trained_outputs = saved_model(images)
      predicted_probability, predicted = torch.max(trained_outputs, 1)

      print("Predicted classes : ", [classes[predicted[idx]] for idx in
      ↪range(batch_size)])
```

```
Predicted classes :  ['dog', 'cat', 'frog', 'horse']
```

```
[22]: from sklearn.metrics import confusion_matrix
      import seaborn as sns

      true_labels_list = []
      predicted_labels_list = []
      model.eval()


      with torch.no_grad():
          for images, labels in testloader:
              images = images.to(device)
              labels = labels.to(device)
              output = model(images)
              predictions = F.softmax(output, dim=1)
              predicted_values, predicted_labels = torch.max(predictions, 1)
              true_labels_list.extend(labels.cpu().numpy())
              predicted_labels_list.extend(predicted_labels.cpu().numpy())

      disp_matrix = confusion_matrix(true_labels_list, predicted_labels_list)
      # Plot the confusion matrix
      plt.figure(figsize=(10, 10))
```

```

sns.heatmap(displacement_matrix, annot=True, fmt="d", cmap="Reds",
            xticklabels=range(10), yticklabels=range(10))
plt.title("Confusion Matrix")
plt.xlabel("Predicted Class Labels")
plt.ylabel("True Class Labels")
plt.show()

```



[22]: